

# Julia Code: Optimized

René Donner  
code@donner.at

1 / 44

What can we do to make our code  
faster?

# "Julia is fast"

Well, it **allows** to get C-like speed

But you can also write very **slow** code

Be aware what you ask the machine to do

# Overview

Don't optimize too early - **profile** first

Know your **memory layout**

**Type stability**

**Compiler hints**

External **tools**

# Profiling

# Profiling

Runtime, allocations, garbage collection

@time

```
function f()  
  for x in 1:100  
    rand(100,1000)  
  end  
end  
@time f() # compiles code  
@time f()
```

```
0.060411 seconds  
(39.85 k allocations: 78.011 MB, 12.62% gc time)  
0.025687 seconds  
(404 allocations: 76.305 MB, 21.90% gc time)
```

# Profiling

@time

```
function f()
    @time r = zeros(2,1)
    for i in 1:3
        @time r = r + rand(2,1)
    end
end
f(); f()
```

```
0.000001 seconds (1 allocation: 80 bytes)
0.000001 seconds (4 allocations: 224 bytes)
0.000000 seconds (4 allocations: 224 bytes)
0.000000 seconds (4 allocations: 224 bytes)
```

# Profiling

@time

```
function f()  
  r = zeros(2,1)  
  for i in 1:3  
    r = r + rand(2,1)  
  end  
end  
f(); @time f()
```

0.000002 seconds (17 allocations: 912 bytes)

```
function f()  
  r = zeros(2,1)  
  for i in 1:3  
    for j in 1:2  
      r[j] += rand()  
    end  
  end  
end  
f(); @time f()
```

0.000001 seconds (5 allocations: 240 bytes)



# Profiling

@time

@profile

```
@profile f()  
Profile.print()  
  
Profile.clear()  
@profile g()  
Profile.print(format = :flat)
```

```
f() = svd(rand(1000,1000))
f(); @profile f()
;
```

```
Profile.print(format = :flat)
```

Count	File	Function	Line
1237	....4/IJulia/src/IJulia.jl	eventloop	141
1	....4/IJulia/src/IJulia.jl	eventloop	162
1236	.../src/execute_request.jl	execute_request_0x535c5df2	177
1	.../src/execute_request.jl	execute_request_0x535c5df2	180
1	.../v0.4/IJulia/src/msg.jl	send_ipython	56
1	.../v0.4/IJulia/src/msg.jl	send_status	112
1	...a/v0.4/JSON/src/JSON.jl	_print	118
1	...a/v0.4/JSON/src/JSON.jl	print	198
628	In[46]	f	1
608	In[47]	f	1
4	arraymath.jl	transpose!	323
4	arraymath.jl	transposeblock!	340
20	arraymath.jl	transposeblock!	346
14	arraymath.jl	transposeblock!	350
22	arraymath.jl	transposeblock!	351
3	dsfmt.jl	dsfmt_fill_array_close_open!	76
1	iostream.jl	sprint	206
2	linalg/lapack.jl	gesdd!	1445
1227	linalg/lapack.jl	gesdd!	1474
1229	linalg/svd.jl	svdfact!	17
1229	linalg/svd.jl	svdfact	23
1236	loading.jl	include_string	266
1236	profile.jl	anonymous	16
3	random.jl	rand!	347
1238	task.jl	anonymous	447

10 / 44

# Profiling

@time

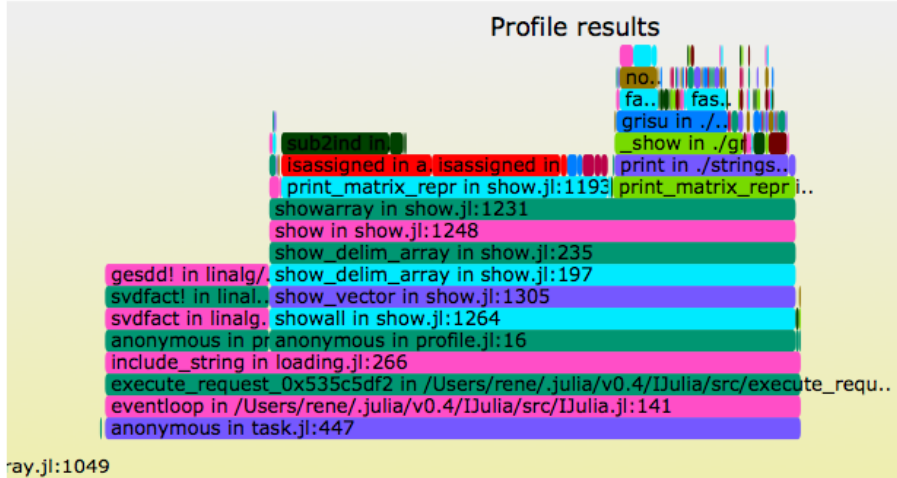
@profile

ProfileView

```
@profile f()  
using ProfileView  
ProfileView.view()
```

```
Profile.clear()
@profile repr([svd(rand(100,i)) for i in rand(100:200,10)])
;
```

```
using ProfileView
ProfileView.view()
```



# Profiling

@time

```
julia --track-allocation=user myscript.jl
```

```
mycode()  
Profile.clear_malloc_data()  
mycode()
```

@profile

Results found in `.mem` files

ProfileView


Memory  
allocations

[http://docs.julialang.org/en/release-0.4/manual/  
profile/#memory-allocation-analysis](http://docs.julialang.org/en/release-0.4/manual/profile/#memory-allocation-analysis)

# Memory Layout

# Memory

## Cache hierarchy



<http://www.1024cores.net/home/parallel-computing/cache-oblivious-algorithms>

## Memory

Row major: First index changes infrequently

### Cache hierarchy

Column major: First index changes most often

### Row major

=> memory layout



# Memory

## Cache hierarchy

## Row major

```
function f(a)
    r = zero(eltype(a))
    for m in 1:size(a,1)
        for n in 1:size(a,2)
            r += a[m,n]
        end
    end
    r
end

function g(a)
    r = zero(eltype(a))
    for n in 1:size(a,2)
        for m in 1:size(a,1)
            r += a[m,n]
        end
    end
    r
end

a = rand(1000,100000)

f(a); @time f(a)
g(a); @time g(a)
```

# Memory

## Cache hierarchy

## Row major

```
function f(a)
    ...
end

function g(a)
    ...
end

a = rand(1000,100000)

f(a); @time f(a)
g(a); @time g(a)
```

0.964728 seconds (5 allocations: 176 bytes)  
0.090653 seconds (5 allocations: 176 bytes)

- Cache aware algorithms
- Cache oblivious algorithms

```

function f(a,x,y)
    for i=1:length(x)
        y[i] += a*x[i]
    end
end

function f_simd(a,x,y)
    @simd for i=1:length(x)
        @inbounds y[i] += a*x[i]
    end
end

n = 1003
x = rand(Float32,n)
y = rand(Float32,n)
f(1.414f0, x, y);

f() = f(1.414f0, x, y);
f_simd() = f_simd(1.414f0, x, y);

```

```

using Benchmark

benchmark(f, "w/o SIMD", 10)
benchmark(f_simd, "with SIMD", 10)

compare([f,f_simd], 100000)

```

	Function	Average	Relative	Replications
1	f	1.31839873e-6	3.195964184707031	100000
2	f_simd	4.125198699999999e-7	1.0	100000

# Type stability

# Type stability

Typed vs  
Any

- Every variable has a type
- Concrete:
  - `Float64`
  - `Int`
- Boxed:
  - `Any`
  - `Union{Int, Float64}`

```

function f(a,b)
    r = 0
    for i = 1:length(a)
        r += a[i] + b[i]
    end
    r
end

function g(a,b)
    r = zero(eltype(a))
    for i = 1:length(a)
        r += a[i] + b[i]
    end
    r
end

a = rand(10_000_000)
b = rand(10_000_000)

f() = f(a,b)
g() = g(a,b)

using Benchmark
benchmark(f, "f", 10); benchmark(g, "g", 10)

compare([f,g], 10)

```

	Function	Average	Relative	Replications
1	f	0.2467122758	16.632886516812558	10
2	g	0.0148327998	1.0	10

# Type stability

Typed vs  
Any

Effect

@code\_warntype

```
@code_warntype f(a,b)
```

```
Variables:
  a::Array{Float64,1}
  b::Array{Float64,1}
  r::ANY
  #s41::Int64
  i::Int64

Body:
  begin # In[4], line 2:
    r = 0 # In[4], line 3:
      GenSym(2) = (Base.arraylen)(a::Array{Float64,1})::Int
      GenSym(0) = $(Expr(:new, UnitRange{Int64}, 1, :(((toF
eld)))(Base.Intrinsics,:select_value)::I)((Base.sle_int)(1,0
2))::Bool,GenSym(2),(Base.box)(Int64,(Base.sub_int)(1,1))::
```

```
@code_warntype g(a,b)
```

```
Variables:
  a::Array{Float64,1}
  b::Array{Float64,1}
  r::Float64
  #s41::Int64
  i::Int64
```

# Compiler hints



## Compiler hints

- `@simd, @inbouds`
- `@inline`
- `@fastmath`

# External tools

# ParallelAccelerator

- From Intel, just released
- Converts Julia to C internally
- Uses OpenMP for parallelization

# ParallelAccelerator

## Unary functions

`-`, `+`, `acos`, `acosh`, `angle`, `asin`, `asinh`, `atan`, `atanh`,  
`cbrt`, `cis`, `cos`, `cosh`, `exp10`, `exp2`, `exp`, `expm1`, `lgamma`,  
`log10`, `log1p`, `log2`, `log`, `sin`, `sinh`, `sqrt`, `tan`, `tanh`,  
`abs`, `copy`, `erf`

## Binary functions

`-`, `+`, `.+`, `.-`, `.*`, `./`, `.\`, `.*`, `.>`, `.<`, `.<=`, `.>=`, `==`,  
`.<`, `.>`, `.^`, `div`, `mod`, `rem`, `&`, `|`, `$`, `min`, `max`

## Parallel Comprehensions

## Stencils

eg. image filtering

# Julia Code: Optimized

René Donner  
code@donner.at

29 / 44



Following: Documentation / examples for  
gistdeck / reveal.js

# How does it work?

-

## Markdown

## - Inside HTML

A simple HTML document is needed for hosting the styles, Markdown and the generated slides themselves:

```
<!DOCTYPE html>
<html>
<head>
<style type="text/css">
/* Slideshow styles */
</style>
</head>
<body>
* <textarea id="source">
<!-- Slideshow Markdown -->
</textarea>
* <script type="text/javascript" src="remark.js">
</script>
<script type="text/javascript">
* var slideshow = remark.create();
</script>
</body>
</html>
```

You may download remark to

32 / 44



Of course, Markdown can only go so far.

# Markdown extensions

To help out with slide layout and formatting, a few Markdown extensions have been included:

- Slide properties, for naming, styling and templating slides
- Content classes, for styling specific content
- Syntax highlighting, supporting a range of languages

# Markdown extensions

## - Slide properties

Initial lines containing key-value pairs are extracted as slide properties:

```
name: agenda
class: middle, center
# Agenda
The name of this slide is {{ name }}.
```

Slide properties serve multiple purposes:

- Naming and styling slides using properties `name` and `class`
- Using slides as templates using properties `template` and `layout`
- Expansion of `{{ property }}`

# Markdown extensions

## - Slide properties

Any occurrences of one or more dotted CSS class names followed by square brackets are replaced with the contents of the brackets with the specified classes applied:

```
.footnote[.red.bold[*] Important footnote]
```

Resulting HTML extract:

```
<span class="footnote">  
<span class="red bold">*</span> Important footnote  
</span>
```

## - Content classes

# Markdown extensions

## - Slide properties

## - Content classes

Code blocks can be syntax highlighted by specifying a language from the set of **supported languages**. Using **GFM** fenced code blocks you can easily specify highlighting language

A number of highlighting **styles** are available, including several well-known themes from different editors and IDEs.

```
```javascript
function add(a, b){
return a + b
```

```
```ruby
def add(a, b)
a + b
```

Markdown extensions  $\int_{\Omega} \nabla v \cdot \nabla u - \lambda v u \, dV = \int_{\Omega} v f \, dV$

- Slide  
properties

-

Content  
classes

## Presenter mode

To help out with giving presentations, a presenter mode comprising the following features is provided:

- Display of slide notes for the current slide, to help you remember key points
- Display of upcoming slide, to let you know what's coming
- Cloning of slideshow for viewing on extended display

# Presenter mode

## - Inline notes

Just like three dashes separate slides, three question marks separate slide content from slide notes:

```
Slide 1 content
*???
Slide 1 notes
---
Slide 2 content
*???
Slide 2 notes
```

Slide notes are also treated as Markdown, and will be converted in the same manner slide content is. Pressing **P** will toggle presenter mode.



# Presenter mode

Presenter mode of course makes no sense to the audience. Creating a cloned view of your slideshow lets you:

## - Inline notes

## - Cloned view

- Move the cloned view to the extended display visible to the audience
- Put the original slideshow in presenter mode
- Navigate as usual, and the cloned view will automatically keep up with the original Pressing **C** will open a cloned view of the current slideshow in a new browser window.

It's time to get started!

# Getting started

Getting up and running is done in only a few steps:

1. Go to your gist and write some markdown.
2. put a "deck" right after "gist" in url like this:

deck  
↓  
<https://gist.github.com/jcouyang/8acfc555a718d62b77b2>  
↓  
<https://gistdeck.github.com/jcouyang/8acfc555a718d62b77b2>

3. enjoy the slide remark generated. For more information on using remark, please check out the [wiki](#) pages.

# That's all folks (for now)!

Slideshow created using [remark](#) and host on [gistdeck](#).

